

Kiosk24

October 26, 2023

Kiosk24

Late-night convenience stores, popularly known as Späti (German trivialization for late-night stores) have a cult status in Germany and have historically (together with Gas-stations) catered the needs for late-night groceries and other essential supplies when normal supermarkets are closed. Kiosk24 is Germany's largest all-day convenience store chain and is the leader in the country's 4 billion Euro late-night convenience store market. Their USP is that they are, unlike regular shops, open round-the-clock and 24/7.

0.0.1 Problem statement

Kiosk24 started a pilot project in Berlin offering bicycle based delivery services. It has been gathering data over a period of one year and wants to get an estimated delivery time that it can provide the customers on the basis of a 10 variable strong feature-set that includes among others the distance from the store to their location, total on-duty delivery partners, no. of items ordered and total outstanding orders.

This dataset hence has the required data to train a regression model that will estimate the delivery-time (dependent variable) based on the feature-set (independent variable).

0.1 Feature Set

store_id: individual id of the store

created_at: the timestamp at which the order was placed

delivered_at: the timestamp when the order was delivered

total_items: total number of items

distinct_items: number of distinct items in the order

subtotal: final bill of the order

total_onshift_runners: no. of delivery partners on duty at order placement

total_busy_runners: number of delivery partners attending to other tasks

total_outstanding_orders: total number of orders to be fulfilled at the moment

store_to_consumer_driving_duration_seconds: approximate travel time from restaurant to customer

0.2 ML Development Lifecycle

0.2.1 Data preperation

importing libraries

```
[39]: import warnings
warnings.filterwarnings("ignore")
```

```
[2]: #for reading and handling the data
import pandas as pd
import numpy as np
import os

#for visualizinng and analyzing data
import matplotlib.pyplot as plt
import seaborn as sns

#data preprocessing
from sklearn.neighbors import LocalOutlierFactor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

#random forest model training
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.ensemble import RandomForestRegressor

#artificial neural network training
from sklearn import preprocessing
from tensorflow.keras import Model
from tensorflow.keras import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import Dense,Dropout,BatchNormalization,LeakyReLU
from sklearn.model_selection import train_test_split
from tensorflow.keras.losses import MeanSquaredLogarithmicError
from tensorflow.keras.losses import MeanSquaredError
from tensorflow.keras.losses import MeanAbsolutePercentageError
from tensorflow.keras.metrics import mean_absolute_percentage_error
from tensorflow.keras.metrics import RootMeanSquaredError
from tensorflow.keras.metrics import MeanAbsoluteError
from tensorflow.keras.optimizers import SGD,Adam
from sklearn.metrics import mean_absolute_percentage_error
from tensorflow.keras.models import load_model

#Modelbit model deployment
import modelbit
```

```
import joblib
```

```
[3]: sns.set(rc={'figure.figsize':(11.7,8.27)})
```

loading data

```
[4]: for dirname, _, filenames in os.walk('C:\Dataset'):  
      for filename in filenames:  
          print(os.path.join(dirname, filename))
```

```
[5]: df=pd.read_excel('Kiosk24_data.xlsx')
```

understanding data as seen by programming environment

```
[6]: df.head()
```

```
[6]:
```

	store_id	created_at	delivered_at	total_items	\
0	1	2015-02-06 22:24:00	2015-02-06 23:11:00	4	
1	2	2015-02-10 21:49:00	2015-02-10 22:33:00	1	
2	2	2015-02-16 00:11:00	2015-02-16 01:06:00	4	
3	1	2015-02-12 03:36:00	2015-02-12 04:35:00	1	
4	1	2015-01-27 02:12:00	2015-01-27 02:58:00	2	

	distinct_items	subtotal	total_onshift_runners	total_busy_runners	\
0	4	39.10	33	14	
1	1	21.59	1	2	
2	3	54.22	8	6	
3	1	17.33	5	6	
4	2	41.14	5	5	

	total_outstanding_orders	store_to_consumer_driving_duration_seconds
0	21	861
1	2	690
2	18	289
3	8	795
4	7	205

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 175777 entries, 0 to 175776
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	store_id	175777 non-null	int64
1	created_at	175777 non-null	datetime64[ns]
2	delivered_at	175777 non-null	datetime64[ns]
3	total_items	175777 non-null	int64
4	distinct_items	175777 non-null	int64

```

5    subtotal                    175777 non-null float64
6    total_onshift_runners      175777 non-null int64
7    total_busy_runners        175777 non-null int64
8    total_outstanding_orders   175777 non-null int64
9    store_to_consumer_driving_duration_seconds 175777 non-null int64
dtypes: datetime64[ns](2), float64(1), int64(7)
memory usage: 13.4 MB

```

0.2.2 Feature engineering

preparing the data for further analysis

```

[8]: df['store_to_consumer_driving_duration_mins']=round(df['store_to_consumer_driving_duration_seconds'], 0)

[9]: df['time_to_delivery']=df['delivered_at'] - df['created_at']

[10]: df['time_to_delivery_mins']=df['time_to_delivery']/pd.Timedelta('60s')

[11]: df['created_at_hour']=df['created_at'].dt.hour
      df['created_at_day']=df['created_at'].dt.dayofweek

[12]: # Define a dictionary to map numeric day of the week to day names
      day_names = {
          0: 'Monday',
          1: 'Tuesday',
          2: 'Wednesday',
          3: 'Thursday',
          4: 'Friday',
          5: 'Saturday',
          6: 'Sunday'
      }

[13]: # Create a new column with day names based on the 'created_at_day' column
      df['created_at_day_name'] = df['created_at_day'].map(day_names)

```

cleaning data

```

[14]: df.
      ↪drop(['store_to_consumer_driving_duration_seconds', 'time_to_delivery', 'created_at', 'delivered_at'])

[15]: df.isna().sum()

[15]: store_id                0
      total_items             0
      distinct_items          0
      subtotal                0
      total_onshift_runners    0
      total_busy_runners       0

```

```

total_outstanding_orders      0
store_to_consumer_driving_duration_mins  0
time_to_delivery_mins         0
created_at_hour               0
created_at_day                0
created_at_day_name           0
dtype: int64

```

```
[16]: df.dropna(inplace=True)
```

```
[17]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 175777 entries, 0 to 175776
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   store_id                             175777 non-null  int64
1   total_items                           175777 non-null  int64
2   distinct_items                        175777 non-null  int64
3   subtotal                              175777 non-null  float64
4   total_onshift_runners                  175777 non-null  int64
5   total_busy_runners                     175777 non-null  int64
6   total_outstanding_orders               175777 non-null  int64
7   store_to_consumer_driving_duration_mins 175777 non-null  float64
8   time_to_delivery_mins                  175777 non-null  float64
9   created_at_hour                        175777 non-null  int64
10  created_at_day                          175777 non-null  int64
11  created_at_day_name                     175777 non-null  object
dtypes: float64(3), int64(8), object(1)
memory usage: 16.1+ MB

```

```
[18]: df['store_id']=df['store_id'].astype('category').cat.codes
```

```
[19]: df['created_at_hour']=df['created_at_hour'].astype('category').cat.codes
```

```
[20]: df['created_at_day']=df['created_at_day'].astype('category').cat.codes
```

```
[21]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 175777 entries, 0 to 175776
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   store_id                             175777 non-null  int8
1   total_items                           175777 non-null  int64
2   distinct_items                        175777 non-null  int64

```

```

3   subtotal                    175777 non-null float64
4   total_onshift_runners       175777 non-null int64
5   total_busy_runners          175777 non-null int64
6   total_outstanding_orders    175777 non-null int64
7   store_to_consumer_driving_duration_mins 175777 non-null float64
8   time_to_delivery_mins       175777 non-null float64
9   created_at_hour             175777 non-null int8
10  created_at_day              175777 non-null int8
11  created_at_day_name         175777 non-null object
dtypes: float64(3), int64(5), int8(3), object(1)
memory usage: 12.6+ MB

```

```
[22]: df.head()
```

```

[22]:   store_id  total_items  distinct_items  subtotal  total_onshift_runners  \
0         0           4             4      39.10                33
1         1           1             1      21.59                 1
2         1           4             3      54.22                 8
3         0           1             1      17.33                 5
4         0           2             2      41.14                 5

      total_busy_runners  total_outstanding_orders  \
0                   14                21
1                   2                 2
2                   6                18
3                   6                 8
4                   5                 7

      store_to_consumer_driving_duration_mins  time_to_delivery_mins  \
0                                   14.0                47.0
1                                   12.0                44.0
2                                   5.0                55.0
3                                   13.0                59.0
4                                   3.0                46.0

      created_at_hour  created_at_day  created_at_day_name
0                   17             4             Friday
1                   16             1             Tuesday
2                    0             0             Monday
3                    3             3             Thursday
4                    2             1             Tuesday


```

```
[23]: #df.to_excel("Kiosk24_data_mod.xlsx", index=False)
```

visualizing data

```
[24]: sns.heatmap(df[[col for col in df.columns if col != 'created_at_day_name']].
    ↪corr())
```

[24]: <Axes: >




```
[25]: plt.scatter(df.time_to_delivery_mins, df.total_outstanding_orders,
    ↪label='dataset', color='blue')

plt.xlabel('time_to_delivery_mins')
plt.ylabel('total_outstanding_orders')
plt.title('Scatter Plot')

# Add a legend to distinguish the two DataFrames
plt.legend()
```


[25]: <matplotlib.legend.Legend at 0x798f07398dc0>



```
[26]: sns.
```

```
↳scatterplot(x='total_items',y='subtotal',hue='distinct_items',palette='colorblind',data=df)
```

```
[26]: <Axes: xlabel='total_items', ylabel='subtotal'>
```

identifying outliers

```
[27]: model1=LocalOutlierFactor()
      #model1.fit(df)
      df['lof_anomaly_score']=model1.fit_predict(df[[col for col in df.columns if col_
      ↪ != 'created_at_day_name']])
```

```
[28]: print("number of outliers : ",(len(df.loc[(df['lof_anomaly_score'] == -1)])))
      df_anomaly=df.loc[(df['lof_anomaly_score'] == -1)]
      df=df.loc[(df['lof_anomaly_score'] == 1)]
```

number of outliers : 801

```
[29]: df.drop(['lof_anomaly_score'],axis=1,inplace=True)
```

<ipython-input-29-a629ff9cb1da>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df.drop(['lof_anomaly_score'],axis=1,inplace=True)

```
[30]: plt.scatter(df.time_to_delivery_mins, df.total_outstanding_orders,
↳ label='Non-anomalous data', color='blue')
plt.scatter(df_anomaly.time_to_delivery_mins, df_anomaly.
↳ total_outstanding_orders, label='Anomalous data', color='red')

plt.xlabel('time_to_delivery_mins')
plt.ylabel('total_outstanding_orders')
plt.title('Scatter Plot')


# Add a legend to distinguish the two DataFrames
plt.legend()
```

[30]: <matplotlib.legend.Legend at 0x798f0ba5a2f0>



```
[31]: sns.scatterplot(x='time_to_delivery_mins',y='total_outstanding_orders',data=df)
```


[31]: <Axes: xlabel='time_to_delivery_mins', ylabel='total_outstanding_orders'>



analyzing data

```
[32]: sns.countplot(x=df.created_at_day)
```


```
[32]: <Axes: xlabel='created_at_day', ylabel='count'>
```



```
[33]: # Specify the order of the x-axis categories
order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday',
        ↪ 'Sunday']
```


```
[34]: sns.countplot(x=df.created_at_day_name, order=order)
```

```
[34]: <Axes: xlabel='created_at_day_name', ylabel='count'>
```



```
[35]: sns.countplot(x=df.created_at_hour)
```

```
[35]: <Axes: xlabel='created_at_hour', ylabel='count'>
```



0.2.3 Data Splitting and Modelling

storing independent and dependent variables seperately

```
[36]: y=df['time_to_delivery_mins']
      x=df.drop(['time_to_delivery_mins', 'created_at_day_name'],axis=1)
```

```
[37]: X_train_RF,X_test_RF,y_train_RF,y_test_RF=train_test_split(x,y,test_size=0.
      ↪2,random_state=42)
```

```
[38]: x.head()
```

```
[38]:
```

	store_id	total_items	distinct_items	subtotal	total_onshift_runners	\
0	0	4	4	39.10	33	
1	1	1	1	21.59	1	
2	1	4	3	54.22	8	
3	0	1	1	17.33	5	
4	0	2	2	41.14	5	

	total_busy_runners	total_outstanding_orders	\
0	14	21	
1	2	2	
2	6	18	

3	6	8
4	5	7

	store_to_consumer_driving_duration_mins	created_at_hour	created_at_day
0	14.0	17	4
1	12.0	16	1
2	5.0	0	0
3	13.0	3	3
4	3.0	2	1

```
[40]: y.head()
```

```
[40]: 0    47.0
      1    44.0
      2    55.0
      3    59.0
      4    46.0
      Name: time_to_delivery_mins, dtype: float64
```

Supervised Learning - Random Forest Regression Model *initializing the RF model*

```
[130]: RF_regressor=RandomForestRegressor()
      RF_regressor.fit(X_train_RF,y_train_RF)
```

```
[130]: RandomForestRegressor()
```

training the RF model

```
[ ]: prediction=RF_regressor.predict(X_test_RF)
```

calculating performance metrics

```
[131]: mse=mean_squared_error(y_test_RF,prediction)
      rmse=mse**.5
      print("mse : ",mse)
      print("rmse : ",rmse)
      mae=mean_absolute_error(y_test_RF,prediction)
      print("mase : ",mae)
```

```
mse : 3.9588216395423164
rmse : 1.9896787779795804
mase : 1.4754132581138526
```

```
[132]: r2_score(y_test_RF,prediction)
```

```
[132]: 0.954812729345795
```

```
[133]: def MAPE(Y_actual,Y_Predicted):
      mape=np.mean(np.abs((Y_actual - Y_Predicted)/Y_actual))*100
```

```
return mape
```


```
[134]: print("mape : ",MAPE(y_test_RF,prediction))
```

```
mape : 3.232319290341489
```

plotting feature-importance metrics

```
[135]: sorted_idx=RF_regressor.feature_importances_.argsort()  
plt.barh(x.columns[sorted_idx],RF_regressor.feature_importances_[sorted_idx])  
plt.xlabel("Random Forest Feature Importance")
```

```
[135]: Text(0.5, 0, 'Random Forest Feature Importance')
```



Supervised Learning - Neural Networks *scaling data before feeding to ANN*

```
[41]: scaler=preprocessing.MinMaxScaler()  
x_scaled=scaler.fit_transform(x)  
X_train_ANN_scaled,X_test_ANN_scaled,y_train_ANN_scaled,y_test_ANN_scaled=train_test_split(x_scaled,  
↪2,random_state=42)
```

configuring the ANN

```
[42]: ANN_model=Sequential()  
ANN_model.add(Dense(14,kernel_initializer='normal',activation='relu'))  
ANN_model.add(Dense(512,activation='relu'))  
ANN_model.add(Dense(1024,activation='relu'))  
ANN_model.add(Dense(256,activation='relu'))
```



```
ANN_model.add(Dense(1,activation='linear'))
```

training the ANN

```
[43]: adam=Adam(learning_rate=0.01)
ANN_model.compile(loss='mse',optimizer=adam,metrics=['mse','mae'])
history=ANN_model.
    ↪fit(X_train_ANN_scaled,y_train_ANN_scaled,epochs=30,batch_size=512,verbose=1,validation_spl
    ↪2)
```

Epoch 1/30

219/219 [=====] - 12s 46ms/step - loss: 217.6809 - mse: 217.6809 - mae: 9.1983 - val_loss: 9.1728 - val_mse: 9.1728 - val_mae: 2.2848

Epoch 2/30

219/219 [=====] - 10s 47ms/step - loss: 6.0582 - mse: 6.0582 - mae: 1.8801 - val_loss: 4.0873 - val_mse: 4.0873 - val_mae: 1.5245

Epoch 3/30

219/219 [=====] - 12s 55ms/step - loss: 4.1660 - mse: 4.1660 - mae: 1.5482 - val_loss: 3.6192 - val_mse: 3.6192 - val_mae: 1.4625

Epoch 4/30

219/219 [=====] - 10s 47ms/step - loss: 3.7729 - mse: 3.7729 - mae: 1.4748 - val_loss: 3.3114 - val_mse: 3.3114 - val_mae: 1.3744

Epoch 5/30

219/219 [=====] - 10s 47ms/step - loss: 3.6301 - mse: 3.6301 - mae: 1.4593 - val_loss: 3.3004 - val_mse: 3.3004 - val_mae: 1.3679

Epoch 6/30

219/219 [=====] - 10s 44ms/step - loss: 3.3700 - mse: 3.3700 - mae: 1.4197 - val_loss: 3.0133 - val_mse: 3.0133 - val_mae: 1.3752

Epoch 7/30

219/219 [=====] - 10s 47ms/step - loss: 2.6348 - mse: 2.6348 - mae: 1.2793 - val_loss: 1.9835 - val_mse: 1.9835 - val_mae: 1.1086

Epoch 8/30

219/219 [=====] - 10s 47ms/step - loss: 2.2882 - mse: 2.2882 - mae: 1.1980 - val_loss: 1.8560 - val_mse: 1.8560 - val_mae: 1.0705

Epoch 9/30

219/219 [=====] - 10s 44ms/step - loss: 2.5618 - mse: 2.5618 - mae: 1.2715 - val_loss: 1.9508 - val_mse: 1.9508 - val_mae: 1.1091

Epoch 10/30

219/219 [=====] - 10s 45ms/step - loss: 2.3258 - mse: 2.3258 - mae: 1.2083 - val_loss: 2.7062 - val_mse: 2.7062 - val_mae: 1.3085

Epoch 11/30

219/219 [=====] - 10s 47ms/step - loss: 2.6713 - mse: 2.6713 - mae: 1.2995 - val_loss: 3.1880 - val_mse: 3.1880 - val_mae: 1.4283

Epoch 12/30

219/219 [=====] - 10s 47ms/step - loss: 2.8852 - mse: 2.8852 - mae: 1.3447 - val_loss: 1.7865 - val_mse: 1.7865 - val_mae: 1.0560

Epoch 13/30


219/219 [=====] - 10s 46ms/step - loss: 2.5955 - mse:

2.5955 - mae: 1.2769 - val_loss: 1.8652 - val_mse: 1.8652 - val_mae: 1.0739
Epoch 14/30
219/219 [=====] - 10s 45ms/step - loss: 2.7296 - mse:
2.7296 - mae: 1.3112 - val_loss: 1.9660 - val_mse: 1.9660 - val_mae: 1.1181
Epoch 15/30
219/219 [=====] - 10s 47ms/step - loss: 2.9996 - mse:
2.9996 - mae: 1.3722 - val_loss: 1.8569 - val_mse: 1.8569 - val_mae: 1.0589
Epoch 16/30
219/219 [=====] - 11s 50ms/step - loss: 2.7881 - mse:
2.7881 - mae: 1.3158 - val_loss: 2.0841 - val_mse: 2.0841 - val_mae: 1.1435
Epoch 17/30
219/219 [=====] - 10s 46ms/step - loss: 2.2762 - mse:
2.2762 - mae: 1.1916 - val_loss: 3.2581 - val_mse: 3.2581 - val_mae: 1.4745
Epoch 18/30
219/219 [=====] - 10s 44ms/step - loss: 2.4843 - mse:
2.4843 - mae: 1.2443 - val_loss: 2.4424 - val_mse: 2.4424 - val_mae: 1.2470
Epoch 19/30
219/219 [=====] - 10s 47ms/step - loss: 2.9288 - mse:
2.9288 - mae: 1.3533 - val_loss: 3.1977 - val_mse: 3.1977 - val_mae: 1.4618
Epoch 20/30
219/219 [=====] - 11s 50ms/step - loss: 2.3910 - mse:
2.3910 - mae: 1.2206 - val_loss: 2.4023 - val_mse: 2.4023 - val_mae: 1.2584
Epoch 21/30
219/219 [=====] - 10s 47ms/step - loss: 2.6477 - mse:
2.6477 - mae: 1.2822 - val_loss: 2.2085 - val_mse: 2.2085 - val_mae: 1.1701
Epoch 22/30
219/219 [=====] - 10s 46ms/step - loss: 2.1466 - mse:
2.1466 - mae: 1.1547 - val_loss: 1.7512 - val_mse: 1.7512 - val_mae: 1.0296
Epoch 23/30
219/219 [=====] - 10s 47ms/step - loss: 2.5557 - mse:
2.5557 - mae: 1.2631 - val_loss: 1.8610 - val_mse: 1.8610 - val_mae: 1.0789
Epoch 24/30
219/219 [=====] - 11s 49ms/step - loss: 2.3334 - mse:
2.3334 - mae: 1.2051 - val_loss: 2.1364 - val_mse: 2.1364 - val_mae: 1.1436
Epoch 25/30
219/219 [=====] - 10s 48ms/step - loss: 2.3079 - mse:
2.3079 - mae: 1.2021 - val_loss: 1.8558 - val_mse: 1.8558 - val_mae: 1.0671
Epoch 26/30
219/219 [=====] - 10s 45ms/step - loss: 2.6191 - mse:
2.6191 - mae: 1.2813 - val_loss: 2.5071 - val_mse: 2.5071 - val_mae: 1.2823
Epoch 27/30
219/219 [=====] - 11s 48ms/step - loss: 2.1046 - mse:
2.1046 - mae: 1.1385 - val_loss: 1.7582 - val_mse: 1.7582 - val_mae: 1.0407
Epoch 28/30
219/219 [=====] - 11s 50ms/step - loss: 2.2636 - mse:
2.2636 - mae: 1.1838 - val_loss: 2.0191 - val_mse: 2.0191 - val_mae: 1.1106
Epoch 29/30
219/219 [=====] - 11s 48ms/step - loss: 2.0492 - mse:

2.0492 - mae: 1.1247 - val_loss: 2.6909 - val_mse: 2.6909 - val_mae: 1.3352
Epoch 30/30
219/219 [=====] - 9s 43ms/step - loss: 2.1712 - mse:
2.1712 - mae: 1.1611 - val_loss: 2.2826 - val_mse: 2.2826 - val_mae: 1.2083

plotting the training history

```
[49]: def plot_history(history, key):  
    plt.plot(history.history[key])  
    plt.plot(history.history['val_'+key])  
    plt.xlabel("Epochs")  
    plt.ylabel(key)  
    plt.legend([key, 'val_'+key])  
    plt.show()  
    #plot the history  
    plot_history(history, 'loss')
```



calculating performance metrics

```
[140]: z_scaled = ANN_model.predict(X_test_ANN_scaled)
```

1094/1094 [=====] - 6s 5ms/step

```
[141]: r2_score(y_test_ANN_scaled, z_scaled)
```

```
[141]: 0.9812296597378213
```

```
[148]: mse = mean_squared_error(y_test_ANN, z)
rmse = mse**.5
print("mse : ",mse)
print("rmse : ",rmse)
print("errors for neural net")
mae = mean_absolute_error(y_test_ANN, z)
print("mae : ",mae)
```

```
mse : 1.8401040600747824
rmse : 1.356504353135213
errors for neural net
mae : 1.0484327619750782
```

```
[149]: mean_absolute_percentage_error(y_test_ANN, z)
```

```
[149]: 0.023490778239000434
```

scaling data before feeding to neural network and creating a custom TensorFlow layer that performs the scaling operation (alernate method)

```
[142]: #
↳X_train_ANN,X_test_ANN,y_train_ANN,y_test_ANN=train_test_split(x,y,test_size=0.
↳2,random_state=42)
```

```
[143]: # import tensorflow as tf
# from sklearn import preprocessing

# class MinMaxScalerLayer(tf.keras.layers.Layer):
#     def __init__(self, scaler_s, **kwargs):
#         super(MinMaxScalerLayer, self).__init__(**kwargs)
#         self.min_ = tf.Variable(initial_value=scaler_s.data_min_,
↳trainable=False, dtype=tf.float32)
#         self.scale_ = tf.Variable(initial_value=scaler_s.scale_,
↳trainable=False, dtype=tf.float32)

#     def call(self, inputs):
#         return (inputs - self.min_) * self.scale_
```

```
[144]: # # Train-test split without scaling
# X_train_ANN, X_test_ANN, y_train_ANN, y_test_ANN = train_test_split(x, y,
↳test_size=0.2, random_state=42)

# # Initialize the MinMaxScaler and fit on the training data
# scaler_scaled = preprocessing.MinMaxScaler()
```

```

# scaler_scaled.fit(X_train_ANN)

# # Create a new sequential model
# ANN_model_scaled = tf.keras.Sequential()

# # Add the scaler layer to the model
# ANN_model_scaled.add(MinMaxScalerLayer(scaler_scaled,
    ↪input_shape=(X_train_ANN.shape[1],)))

# # Rest of your model architecture
# ANN_model_scaled.add(Dense(14, kernel_initializer='normal',
    ↪activation='relu'))
# ANN_model_scaled.add(Dense(512, activation='relu'))
# ANN_model_scaled.add(Dense(1024, activation='relu'))
# ANN_model_scaled.add(Dense(256, activation='relu'))
# ANN_model_scaled.add(Dense(1, activation='linear'))

# # Compile and train the model
# adam = Adam(learning_rate=0.01)
# ANN_model_scaled.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae'])
# history = ANN_model_scaled.fit(X_train_ANN, y_train_ANN, epochs=30,
    ↪batch_size=512, verbose=1, validation_split=0.2)

```

Epoch 1/30

219/219 [=====] - 19s 82ms/step - loss: 77.3646 - mse: 77.3646 - mae: 5.0478 - val_loss: 6.1211 - val_mse: 6.1211 - val_mae: 1.9568

Epoch 2/30

219/219 [=====] - 17s 79ms/step - loss: 3.7414 - mse: 3.7414 - mae: 1.5002 - val_loss: 3.8366 - val_mse: 3.8366 - val_mae: 1.6054

Epoch 3/30

219/219 [=====] - 18s 80ms/step - loss: 2.8101 - mse: 2.8101 - mae: 1.3280 - val_loss: 4.9620 - val_mse: 4.9620 - val_mae: 1.9026

Epoch 4/30

219/219 [=====] - 17s 79ms/step - loss: 2.6207 - mse: 2.6207 - mae: 1.2849 - val_loss: 3.8538 - val_mse: 3.8538 - val_mae: 1.6082

Epoch 5/30

219/219 [=====] - 18s 81ms/step - loss: 2.7313 - mse: 2.7313 - mae: 1.3157 - val_loss: 2.6951 - val_mse: 2.6951 - val_mae: 1.3363

Epoch 6/30

219/219 [=====] - 18s 82ms/step - loss: 2.7227 - mse: 2.7227 - mae: 1.3083 - val_loss: 2.0646 - val_mse: 2.0646 - val_mae: 1.1295

Epoch 7/30

219/219 [=====] - 17s 78ms/step - loss: 2.9241 - mse: 2.9241 - mae: 1.3666 - val_loss: 1.8462 - val_mse: 1.8462 - val_mae: 1.0708

Epoch 8/30

219/219 [=====] - 17s 79ms/step - loss: 2.5818 - mse: 2.5818 - mae: 1.2747 - val_loss: 2.5377 - val_mse: 2.5377 - val_mae: 1.2885

Epoch 9/30
 219/219 [=====] - 18s 82ms/step - loss: 2.7673 - mse: 2.7673 - mae: 1.3195 - val_loss: 2.1464 - val_mse: 2.1464 - val_mae: 1.1560
 Epoch 10/30
 219/219 [=====] - 18s 80ms/step - loss: 2.4600 - mse: 2.4600 - mae: 1.2420 - val_loss: 3.2136 - val_mse: 3.2136 - val_mae: 1.4381
 Epoch 11/30
 219/219 [=====] - 17s 77ms/step - loss: 2.5795 - mse: 2.5795 - mae: 1.2662 - val_loss: 1.7182 - val_mse: 1.7182 - val_mae: 1.0252
 Epoch 12/30
 219/219 [=====] - 18s 81ms/step - loss: 2.4844 - mse: 2.4844 - mae: 1.2454 - val_loss: 1.8483 - val_mse: 1.8483 - val_mae: 1.0699
 Epoch 13/30
 219/219 [=====] - 17s 76ms/step - loss: 2.5660 - mse: 2.5660 - mae: 1.2667 - val_loss: 3.8670 - val_mse: 3.8670 - val_mae: 1.6364
 Epoch 14/30
 219/219 [=====] - 18s 80ms/step - loss: 2.2841 - mse: 2.2841 - mae: 1.1859 - val_loss: 1.6720 - val_mse: 1.6720 - val_mae: 1.0112
 Epoch 15/30
 219/219 [=====] - 18s 83ms/step - loss: 2.7089 - mse: 2.7089 - mae: 1.3020 - val_loss: 4.1910 - val_mse: 4.1910 - val_mae: 1.7212
 Epoch 16/30
 219/219 [=====] - 18s 82ms/step - loss: 2.5793 - mse: 2.5793 - mae: 1.2701 - val_loss: 1.7064 - val_mse: 1.7064 - val_mae: 1.0179
 Epoch 17/30
 219/219 [=====] - 18s 83ms/step - loss: 2.3666 - mse: 2.3666 - mae: 1.2127 - val_loss: 1.9628 - val_mse: 1.9628 - val_mae: 1.1205
 Epoch 18/30
 219/219 [=====] - 19s 87ms/step - loss: 2.2150 - mse: 2.2150 - mae: 1.1738 - val_loss: 1.6524 - val_mse: 1.6524 - val_mae: 0.9973
 Epoch 19/30
 219/219 [=====] - 19s 87ms/step - loss: 2.3610 - mse: 2.3610 - mae: 1.2032 - val_loss: 1.6513 - val_mse: 1.6513 - val_mae: 0.9987
 Epoch 20/30
 219/219 [=====] - 19s 87ms/step - loss: 2.2194 - mse: 2.2194 - mae: 1.1714 - val_loss: 6.9084 - val_mse: 6.9084 - val_mae: 2.3159
 Epoch 21/30
 219/219 [=====] - 18s 80ms/step - loss: 2.1262 - mse: 2.1262 - mae: 1.1443 - val_loss: 1.8176 - val_mse: 1.8176 - val_mae: 1.0637
 Epoch 22/30
 219/219 [=====] - 18s 81ms/step - loss: 1.9833 - mse: 1.9833 - mae: 1.1047 - val_loss: 2.6490 - val_mse: 2.6490 - val_mae: 1.3082
 Epoch 23/30
 219/219 [=====] - 17s 79ms/step - loss: 2.2862 - mse: 2.2862 - mae: 1.1944 - val_loss: 1.7825 - val_mse: 1.7825 - val_mae: 1.0515
 Epoch 24/30
 219/219 [=====] - 17s 78ms/step - loss: 2.1635 - mse: 2.1635 - mae: 1.1579 - val_loss: 1.6716 - val_mse: 1.6716 - val_mae: 1.0095

```


Epoch 25/30
219/219 [=====] - 17s 80ms/step - loss: 2.0474 - mse:
2.0474 - mae: 1.1278 - val_loss: 1.7312 - val_mse: 1.7312 - val_mae: 1.0296
Epoch 26/30
219/219 [=====] - 18s 82ms/step - loss: 2.0860 - mse:
2.0860 - mae: 1.1397 - val_loss: 1.6237 - val_mse: 1.6237 - val_mae: 0.9887
Epoch 27/30
219/219 [=====] - 17s 79ms/step - loss: 1.9998 - mse:
1.9998 - mae: 1.1122 - val_loss: 1.6185 - val_mse: 1.6185 - val_mae: 0.9900
Epoch 28/30
219/219 [=====] - 18s 81ms/step - loss: 2.2873 - mse:
2.2873 - mae: 1.1865 - val_loss: 1.6676 - val_mse: 1.6676 - val_mae: 0.9992
Epoch 29/30
219/219 [=====] - 18s 80ms/step - loss: 1.8838 - mse:
1.8838 - mae: 1.0777 - val_loss: 1.7432 - val_mse: 1.7432 - val_mae: 1.0459
Epoch 30/30
219/219 [=====] - 17s 80ms/step - loss: 2.1644 - mse:
2.1644 - mae: 1.1580 - val_loss: 1.7986 - val_mse: 1.7986 - val_mae: 1.0381

```

```

[145]: # def plot_history(history, key):
#       plt.plot(history.history[key])
#       plt.plot(history.history['val_'+key])
#       plt.xlabel("Epochs")
#       plt.ylabel(key)
#       plt.legend([key, 'val_'+key])
#       plt.show()
# #plot the history
# plot_history(history, 'loss')

```



```
[146]: # z=ANN_model_scaled.predict(X_test_ANN)
```

```
1094/1094 [=====] - 5s 5ms/step
```

```
[147]: # r2_score(y_test_ANN, z)
```

```
[147]: 0.9789964570861258
```

0.2.4 Model Evaluation

saving the model

```
[150]: joblib.dump(RF_regressor, 'RF.joblib')
```

```
[150]: ['RF.joblib']
```

```
[45]: joblib.dump(scaler, "scaler.joblib")
```

```
[45]: ['scaler.joblib']
```

```
[46]: # Saving the model in SavedModel format
ANN_model.save("ANN")
```



```
[155]: # Saving the model in SavedModel format
# ANN_model_scaled.save("ANN_scaled")
```

loading the model saved before

```
[217]: # Load the model
RF_loaded_model = joblib.load("RF.joblib")
```

```
[47]: # Load the model from the SavedModel format
ANN_loaded_model = load_model("ANN")
# ANN_loaded_model_scaled = load_model("ANN_scaled")
#options = LoadOptions(experimental_io_device="/job:localhost")
#ANN_loaded_model = tf.keras.models.load_model("ANN", options=options)

# Load the MinMaxScaler using joblib
loaded_scaler = joblib.load("scaler.joblib")
```

testing the model

```
[48]: ANN_loaded_model.predict(loaded_scaler.transform(x.iloc[0:1])).flatten()
```

```
1/1 [=====] - 0s 71ms/step
```

```
[48]: array([45.296677], dtype=float32)
```

```
[159]: # ANN_loaded_model_scaled.predict(x.iloc[0:1]).flatten()
```

```
1/1 [=====] - 0s 123ms/step
```

```
[159]: array([46.29598], dtype=float32)
```

defining predictor functions

```
[230]: def RF_regressor_predictor(sample):

    """
    Predict the target value for a single sample using a regression model.
    """

    # Extract and return the predicted value
    return RF_loaded_model.predict(sample)
```

```
[206]: def RF_regressor_predictor_2(sample):

    """
    Predict the target value for a single sample using a regression model.
    """

    # Extract and return the predicted value
    return RF_regressor.predict(sample)
```

```
[105]: def ANN_regressor_predictor(sample):

        """
        Predict the target value for a single sample using a regression model.
        """

        # Extract and return the predicted value
        return ANN_loaded_model.predict(loaded_scaler.transform(sample)).flatten()
```

```
[104]: def ANN_regressor_predictor_2(sample):

        """
        Predict the target value for a single sample using a regression model.
        """

        # Extract and return the predicted value
        return ANN_model.predict(loaded_scaler.transform(sample)).flatten()
```

```
[163]: def ANN_regressor_predictor_scaled(sample):

        """
        Predict the target value for a single sample using a regression model.
        """

        # Extract and return the predicted value
        return ANN_loaded_model_scaled.predict(sample).flatten()
```

testing the predictor functions

```
[52]: example_input = x.iloc[0:1]
```

```
[220]: RF_regressor_predictor(example_input)
```

```
[220]: array([46.39])
```

```
[208]: RF_regressor_predictor_2(example_input)
```

```
[208]: array([46.39])
```

```
[106]: ANN_regressor_predictor(example_input)
```

```
1/1 [=====] - 0s 17ms/step
```

```
[106]: array([45.296677], dtype=float32)
```

```
[107]: ANN_regressor_predictor_2(example_input)
```

```
1/1 [=====] - 0s 20ms/step
```

```
[107]: array([45.296677], dtype=float32)
```

```
[169]: ANN_regressor_predictor_scaled(example_input)
```

```
1/1 [=====] - 0s 34ms/step
```

```
[169]: array([46.29598], dtype=float32)
```

0.2.5 Model deployment

deploying API

```
[62]: mb = modelbit.login()
```

```
<IPython.core.display.HTML object>
```

```
[231]: mb.deploy(RF_regressor_predictor,
               dataframe_mode=True,
               example_dataframe=example_input, extra_files = "RF.joblib",
               ↪python_packages=["scikit-learn==1.2.2"])
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
[ ]: mb.deploy(RF_regressor_predictor_2,
               dataframe_mode=True,
               example_dataframe=example_input, python_packages=["scikit-learn==1.2.
               ↪2"])
```

```
[108]: mb.deploy(ANN_regressor_predictor,
               dataframe_mode=True,
               example_dataframe=example_input, python_packages=["tensorflow==2.14.
               ↪0"])
```

```
<IPython.core.display.HTML object>
```

```
Uploading 'ANN_loaded_model': 100%|          | 2.93M/2.93M [00:00<00:00,
4.05MB/s]
```

```
<IPython.core.display.HTML object>
```

```
[109]: mb.deploy(ANN_regressor_predictor_2,
               dataframe_mode=True,
               example_dataframe=example_input, python_packages=["tensorflow==2.14.
               ↪0", "numpy==1.23.5"])
```

```
<IPython.core.display.HTML object>
```

```
Uploading 'ANN_model': 100%|          | 7.79M/7.79M [00:00<00:00, 9.31MB/s]
```

```
<IPython.core.display.HTML object>
```

testing the API

```
[232]: # API on modelbit
import requests, json

requests.post(
    "https://vishwarathtomar.app.modelbit.com/v1/RF_regressor_predictor/latest",
    headers={"Content-Type": "application/json"},
    data=json.dumps({"data": {
        "store_id": 0,
        "total_items": 4,
        "distinct_items": 4,
        "subtotal": 39.1,
        "total_onshift_runners": 33,
        "total_busy_runners": 14,
        "total_outstanding_orders": 21,
        "store_to_consumer_driving_duration_mins": 14,
        "created_at_hour": 17,
        "created_at_day": 4
    }})).json()
```

```
[232]: {'data': 46.39}
```

```
[216]: # API on modelbit
import requests, json

requests.post(
    "https://vishwarathtomar.app.modelbit.com/v1/RF_regressor_predictor_2/
    ↪latest",
    headers={"Content-Type": "application/json"},
    data=json.dumps({"data": {
        "store_id": 0,
        "total_items": 4,
        "distinct_items": 4,
        "subtotal": 39.1,
        "total_onshift_runners": 33,
        "total_busy_runners": 14,
        "total_outstanding_orders": 21,
        "store_to_consumer_driving_duration_mins": 14,
        "created_at_hour": 17,
        "created_at_day": 4
    }})).json()
```

```
[216]: {'data': 46.39}
```

```
[111]: # API on modelbit
import requests, json
```

```
requests.post(
    "https://vishwarathtomar.app.modelbit.com/v1/ANN_regressor_predictor/
↳latest",
    headers={"Content-Type":"application/json"},
    data=json.dumps({"data": {
        "store_id": 0,
        "total_items": 4,
        "distinct_items": 4,
        "subtotal": 39.1,
        "total_onshift_runners": 33,
        "total_busy_runners": 14,
        "total_outstanding_orders":21,
        "store_to_consumer_driving_duration_mins": 14,
        "created_at_hour": 17,
        "created_at_day": 4
    }}}).json()
```

[111]: {'data': 45.29667663574219}

```
[110]: # API on modelbit
import requests, json

requests.post(
    "https://vishwarathtomar.app.modelbit.com/v1/ANN_regressor_predictor_2/
↳latest",
    headers={"Content-Type":"application/json"},
    data=json.dumps({"data": {
        "store_id": 0,
        "total_items": 4,
        "distinct_items": 4,
        "subtotal": 39.1,
        "total_onshift_runners": 33,
        "total_busy_runners": 14,
        "total_outstanding_orders":21,
        "store_to_consumer_driving_duration_mins": 14,
        "created_at_hour": 17,
        "created_at_day": 4
    }}}).json()
```

[110]: {'data': 45.29667663574219}

```
[ ]: # API on local host
import requests, json

requests.post(
    "http://localhost:8605/v1/models/my_model:predict",
    headers={"Content-Type":"application/json"},
```

```
data=json.dumps({"data": {  
    "store_id": 0,  
    "total_items": 4,  
    "distinct_items": 4,  
    "subtotal": 39.1,  
    "total_onshift_runners": 33,  
    "total_busy_runners": 14,  
    "total_outstanding_orders":21,  
    "store_to_consumer_driving_duration_mins": 14,  
    "created_at_hour": 17,  
    "created_at_day": 4  
}})).json()
```

```
[1]: !export PATH=/Library/TeX/texbin:$PATH
```

Der Befehl "export" ist entweder falsch geschrieben oder
konnte nicht gefunden werden.

```
[ ]:
```